

CODEX + SGLANG / ASCEND - Daily Rules

Default: open ~/code/sglang with VS Code in WSL and use the Codex sidebar. Repo rules/skills handle the workflow.

START HERE - 4 QUESTIONS

1. **Same goal + same strategy?** -> continue.
2. **Same goal + new strategy?** -> NEW SESSION.
3. **New objective?** -> NEW SESSION.
4. **Parallel work?** -> another git worktree; one active Codex session per worktree.

LARGE LOG - MUST REDUCE FIRST

Trigger: ≥ 1 MiB OR $\geq 10,000$ lines.

size check -> `extract-log-context.py` ->
`.codex-artifacts/logs/*.focused.txt` -> inspect focused artifact -> narrow raw ranges only if needed.

Never read/cat the whole raw log first. You do not need to tell Codex to use the reducer.

SEARCH ROUTING

Exact symbol/error/path -> **rg**.
Known PR/commit -> **git show/diff**.
Model history -> **model-pr-history**.
Unknown concept -> **Semble**.
Known symbol relations -> **Serena** (optional).

Once a symbol is found, stop broad semantic discovery.

HANDOFF AUTO-DISCOVERY

Review/investigation -> `.codex-artifacts/handoffs/...` -> STOP.

New implementation session -> **SessionStart** hook resolves **PR/branch/worktree** -> injects handoff path -> read -> re-verify -> patch -> consume.

Normally prompt only: **Address the review findings**.
Do not redo the broad review.

SESSION + MODEL ROUTING

Luna/low: metadata, PR description, docs, mechanical.
Terra/medium: default development, conflicts, comments, normal bugs/refactors.
Sol/high: deep review, hard Ascend correctness, HCCL/NPUGraph/perf root cause.
Sol/xhigh: escalation only.

Long perf/kernel loop -> `/plan + /goal + artifact ledger`.

DO NOT SPEND A MODEL TURN

Standalone push/status/fetch/switch -> terminal.
Strategy reversal (instead / undo / restore) -> **NEW SESSION**, not another turn in a long thread.
Do not expand scope after the objective is solved.

ASCEND / NPU HARD GATES

Record before version-sensitive diagnosis/benchmark:

hardware | CANN | torch | torch_npu | sgl-kernel-npu | SGLang commit | graph/eager | dtype/quant | TP/DP/EP | workload | actual backend.

STOP if baseline/candidate differ by unrelated variables or a silent fallback changes backend.

CUDA/NCCL/graph/stream assumptions are not automatically valid for Ascend/HCCL/NPUGraph.

STOP RULES

Before ending: objective complete? diff minimal? validation sufficient? handoff/Goal updated if needed?

No unrelated cleanup/docs/tests/refactors after the objective is solved.

Tests only when explicitly requested, required by review/CI, or needed as a real regression guard.

DAILY UI / QUICK COMMANDS

Default: `cd ~/code/sglang && code .` -> WSL: Ubuntu -> Codex sidebar.

CLI optional: `cxl Luna | cx Terra | cxh Sol/high | cxx Sol/xhigh`

Handoff diagnostic: `python3 .codex/scripts/resolve-handoff.py --json`

Static: `.codex/` | Writable: `.codex-artifacts/`