

CODEX + SGLANG / ASCEND - Daily Rules

Describe the engineering objective. Repo rules/skills handle log reduction, handoffs, and search routing.

START HERE - 4 QUESTIONS

1. **New objective?** -> new session.
2. **Review found actionable local issues?** -> handoff before STOP, then new implementation session.
3. **Large local/downloaded log?** -> reducer before any full raw-log read.
4. **Just Git plumbing?** push/status/fetch/switch -> terminal.

LARGE LOG - MUST REDUCE FIRST

Trigger: ≥ 1 MiB OR $\geq 10,000$ lines.

size check -> `extract-log-context.py` ->
`.codex/logs/*.focused.txt` -> inspect focused artifact -> narrow raw ranges only if needed.

Never read/cat the whole raw log first. You do not need to tell Codex to use the reducer.

SEARCH ROUTING

Exact symbol/error/path -> **rg**.
Known PR/commit -> **git show/diff**.
Model history -> **model-pr-history**.
Unknown concept -> **Semble**.
Known symbol relations -> **Serena** (optional).

Once a symbol is found, stop broad semantic discovery.

REVIEW -> HANDOFF -> NEW SESSION

If report-only review/investigation finds actionable issues and GitHub review threads are not already authoritative:

review -> `.codex/handoffs/. . .` -> STOP -> new implementation session -> re-verify -> patch.

Implementation **does not redo the broad review**.

Handoff only stores: severity | file/symbol | root cause | intended change | constraints | validation.

SESSION + MODEL ROUTING

Luna/low: metadata, PR description, docs, mechanical.

Terra/medium: default development, conflicts, comments, normal bugs/refactors.

Sol/high: deep review, hard Ascend correctness, HCCL/NPUGraph/perf root cause.

Sol/xhigh: escalation only.

Long perf/kernel loop -> `/plan + /goal + artifact ledger`.

DO NOT SPEND A MODEL TURN

Standalone push/status/fetch/switch -> terminal.
Do not send push, continue, or update in a long thread.
Do not expand scope after the objective is solved.

ASCEND / NPU HARD GATES

Record before version-sensitive diagnosis/benchmark:

hardware | CANN | torch | torch_npu | sgl-kernel-npu | SGLang commit | graph/eager | dtype/quant | TP/DP/EP | workload | actual backend.

STOP if baseline/candidate differ by unrelated variables or a silent fallback changes backend.

CUDA/NCCL/graph/stream assumptions are not automatically valid for Ascend/HCCL/NPUGraph.

STOP RULES

Before ending: objective complete? diff minimal? validation sufficient? handoff/Goal updated if needed?

No unrelated cleanup/docs/tests/refactors after the objective is solved.

Tests only when explicitly requested, required by review/CI, or needed as a real regression guard.

QUICK COMMANDS

`cxl Luna | cx Terra | cxh Sol/high | cxx Sol/xhigh`

`/review /side /compact /goal /usage /status`

Handoff: `.codex/scripts/new-handoff.sh <PR|slug>`