

SGLANG + ASCEND - Development Daily Rules

PRIMARY PRIORITY: correct, reproducible SGLang + Ascend engineering. Model routing may reduce cost, never weaken NPU evidence or validation gates.

START HERE - 4 QUESTIONS

1. **Same goal + same strategy?** -> continue.
2. **Same goal + new strategy?** -> NEW SESSION.
3. **New objective?** -> NEW SESSION.
4. **Parallel work?** -> another git worktree; one editing agent per worktree.

LARGE LOG - MUST REDUCE FIRST

Trigger: `>= 1 MiB` OR `>= 10,000` lines.

size check -> `extract-log-context.py` -> `.codex-artifacts/logs/*.focused.txt` -> inspect focused artifact -> narrow raw ranges only if needed.

Never read/cat the whole raw log first. The workflow rule applies regardless of the selected model backend.

SEARCH ROUTING

Exact symbol/error/path -> **rg**.
Known PR/commit -> `git show/diff`.
Model history -> **model-pr-history**.
Unknown concept -> **Semble**.
Known symbol relations -> **Serena** (optional).

Once a symbol is found, stop broad semantic discovery.

HANDOFF AUTO-DISCOVERY

Review/investigation -> `.codex-artifacts/handoffs/...` -> STOP.

New implementation session -> resolve PR/branch/worktree -> read matching handoff -> re-verify -> patch -> consume.

Codex can inject the handoff automatically through hooks; other harnesses run the resolver preflight. **Do not redo the broad review.**

SESSION + MODEL ROUTING

primary: every task uses the selected backend by default.
local / cheap / strong: optional OpenAI-compatible worker slots.
hybrid: opt-in task-class routing.
Hard SGLang/Ascend classes: selected PRIMARY first; workers only after it.

Inspect: `ai-task --dry-run ...` Explicit tier wrappers override routing, not engineering gates.

DO NOT SPEND A MODEL TURN

Standalone push/status/fetch/switch -> terminal.
Strategy reversal (`instead / undo / restore`) -> **NEW SESSION**, not another turn in a long thread.
Do not expand scope after the objective is solved.

ASCEND / NPU HARD GATES

Record before version-sensitive diagnosis/benchmark:

hardware | CANN | torch | torch_npu | sgl-kernel-npu | SGLang commit | graph/eager | dtype/quant | TP/DP/EP | workload | actual backend.

STOP if baseline/candidate differ by unrelated variables or a silent fallback changes backend.

CUDA/NCCL/graph/stream assumptions are not automatically valid for Ascend/HCC/CL/NPUGraph.

STOP RULES

Before ending: objective complete? diff minimal? validation sufficient? handoff/Goal updated if needed?

No unrelated cleanup/docs/tests/refactors after the objective is solved.

Tests only when explicitly requested, required by review/CI, or needed as a real regression guard.

DAILY UI / QUICK COMMANDS

Workspace: `cd ~/code/sglang && code .` -> WSL: Ubuntu -> selected model client.

Router: `ai-task docs|bug|review|npu ...`
Overrides: `local-task | cheap-task | strong-task`; Codex shortcuts exist only when installed.

Handoff diagnostic: `python3 .codex/scripts/resolve-handoff.py --json`

Static: `.codex/` | Writable: `.codex-artifacts/`