

SGLang + Ascend - Daily Workflow Rules

Codex/OpenAI default UI; GLM optional provider; correctness gates are provider-independent.

SESSION

New objective -> new session
New strategy -> new session
One editing agent per worktree

PROVIDERS

OpenAI remains default Codex config
GLM uses isolated low/high/max profiles
balanced: hard NPU/perf/kernel -> OpenAI

ASCEND BASELINE

hardware + CANN + torch + torch_npu
SGLang/sgl-kernel-npu commits
precision + TP/DP/EP + graph + workload

LARGE LOG

>=1 MiB or >=10,000 lines
Run extract-log-context.py first
Read narrow raw ranges only if needed

HANDOFFS

Implementation resolves handoff first
Re-verify against current HEAD
Consume only after all items are done

PERFORMANCE

One hypothesis + one scoped change
Correctness + identical benchmark
Microbenchmark != E2E serving result